

▼ Data loading

▼ Import

```
!pip install kaggle
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.6/dist-packages (1.5.10)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.6/dist-packages (from kaggle) (1.24.3)
Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from kaggle) (2.23.0)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.6/dist-packages (from kaggle) (1.15.0)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.6/dist-packages (from kaggle) (4.0.1)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.6/dist-packages (from kaggle) (2.8.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.6/dist-packages (from kaggle) (4.41.1)
Requirement already satisfied: certifi in /usr/local/lib/python3.6/dist-packages (from kaggle) (2020.12.5)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests->kaggle) (3.0.2)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests->kaggle) (2.10)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.6/dist-packages (from python-slugify->kaggle) (1.3)
```

```
pip install --upgrade kaggle
```

```
Requirement already up-to-date: kaggle in /usr/local/lib/python3.6/dist-packages (1.5.10)
Requirement already satisfied, skipping upgrade: certifi in /usr/local/lib/python3.6/dist-packages (from kaggle) (2020.12.5)
Requirement already satisfied, skipping upgrade: six>=1.10 in /usr/local/lib/python3.6/dist-packages (from kaggle) (1.15.0)
Requirement already satisfied, skipping upgrade: tqdm in /usr/local/lib/python3.6/dist-packages (from kaggle) (4.41.1)
Requirement already satisfied, skipping upgrade: python-slugify in /usr/local/lib/python3.6/dist-packages (from kaggle) (4.0.1)
Requirement already satisfied, skipping upgrade: requests in /usr/local/lib/python3.6/dist-packages (from kaggle) (2.23.0)
Requirement already satisfied, skipping upgrade: python-dateutil in /usr/local/lib/python3.6/dist-packages (from kaggle) (2.8.1)
Requirement already satisfied, skipping upgrade: urllib3 in /usr/local/lib/python3.6/dist-packages (from kaggle) (1.24.3)
Requirement already satisfied, skipping upgrade: text-unidecode>=1.3 in /usr/local/lib/python3.6/dist-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied, skipping upgrade: idna<3,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests->kaggle) (2.10)
Requirement already satisfied, skipping upgrade: chardet<4,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests->kaggle) (3.0.2)
```

```
!mkdir .kaggle
```

```
from google.colab import drive
drive.mount('/content/gdrive')
```

```
Mounted at /content/gdrive
```

```
import os
os.environ['KAGGLE_CONFIG_DIR'] = "/content/gdrive/My Drive/Study/Kaggle"
```

```
%cd /content/gdrive/My Drive/Study/DL/Assignments/Project
```

```
/content/gdrive/My Drive/Study/DL/Assignments/Project
```

```
!kaggle competitions download -c final-project-dl-spring-2020
```

```
Warning: Looks like you're using an outdated API Version, please consider updating (server 1.5.10 / client 1.5.4)
404 - Not Found
```

```
!ls
```

```
alexnet_pretrained.pt      preds_4.1.csv
final-project-dl-spring-2020.zip  preds_4.1.gsheet
model_5.pt                  preds_4.2.csv
model_6.pt                  preds_4.3.csv
model_7.pt                  preds_5.csv
model_8.pt                  preds_6.csv
```

pred_5.csv	preds_7.csv
preds_1.2.csv	preds_7.gsheet
preds_1.2.gsheet	preds_8.csv
preds_1.3.csv	preds_8.gsheet
preds_1.4.csv	Project.docx
preds_1.6.csv	resnet18_pretrained.pt
preds_1.6.gsheet	resnet34_pretrained.pt
'preds_1.7 (1).gsheet'	resnet50_pretrained.pt
preds_1.7.csv	'~\$project.docx'
preds_1.7.gsheet	sample_submission.csv
preds_1.csv	submission.csv
preds_1.gsheet	test_data
preds_3.1.csv	train_data
preds_3.1.gsheet	

```
#unzipping the zip files and deleting the zip files
#!unzip *.zip && rm *.zip
```

```
import torch
import torch.nn as nn

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
from datetime import datetime
from pathlib import Path
import pandas as pd
import torch.nn.functional as F
import torchtext.data as ttd
import torchvision
from torchvision import datasets, transforms, utils

from torch.optim.lr_scheduler import StepLR, CyclicLR
```

```
folder=Path('/content/gdrive/My Drive/Study/DL/Assignments/Project')
```

▼ Data transformation

```
transformer = transforms.Compose([
#   transforms.Resize(256),
#   transforms.RandomCrop(256),
#   transforms.RandomHorizontalFlip(),
#   transforms.RandomVerticalFlip(),
    transforms.ToTensor(),
#   transforms.Normalize([meanR, meanG, meanB], [stdR, stdG, stdB]),
])
```

```
train0_ds = datasets.ImageFolder(folder/'train_data/train',transform=transformer)
test_ds = datasets.ImageFolder(folder/'test_data',transform=transformer)
```

```
train0_ds.classes

['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
```

▼ Train/val split

```
from sklearn.model_selection import StratifiedShuffleSplit

sss = StratifiedShuffleSplit(n_splits=1, train_size=0.7, random_state=0)
```

```

indices=list(range(len(train0_ds)))
y_train0=[y for _,y in train0_ds]
for train_index, val_index in sss.split(indices, y_train0):
    print("train:", train_index, "val:", val_index)
    print(len(val_index),len(train_index))

    train: [5279 5574 7391 ... 8774 2047 3129] val: [1900 5887 8619 ... 5449 1106 6628]
    2708 6317

```

```

from torch.utils.data import Subset

val_ds=Subset(train0_ds,val_index)
train_ds=Subset(train0_ds,train_index)

```

```

print(f'Number of training examples: {len(train_ds)}')
print(f'Number of validation examples: {len(val_ds)}')
print(f'Number of testing examples: {len(test_ds)}')

```

```

Number of training examples: 6317
Number of validation examples: 2708
Number of testing examples: 3929

```

```

import numpy as np

# RGB mean and std
meanRGB=[np.mean(x.numpy(),axis=(1,2)) for x,_ in train_ds]
stdRGB=[np.std(x.numpy(),axis=(1,2)) for x,_ in train_ds]

meanR=np.mean([m[0] for m in meanRGB])
meanG=np.mean([m[1] for m in meanRGB])
meanB=np.mean([m[2] for m in meanRGB])

stdR=np.mean([s[0] for s in stdRGB])
stdG=np.mean([s[1] for s in stdRGB])
stdB=np.mean([s[2] for s in stdRGB])

print(meanR,meanG,meanB)
print(stdR,stdG,stdB)

0.48499483 0.45491496 0.39288142
0.22034366 0.21472108 0.21670675

```

▼ Preprocessing

```

#additional transformation changes if required

train_transformer_1 = transforms.Compose([
    transforms.Resize(512),
    # transforms.RandomHorizontalFlip(), # randomly flip and rotate
    transforms.RandomCrop(512),
    # transforms.RandomHorizontalFlip(),
    # transforms.RandomVerticalFlip(),
    transforms.ToTensor(),
    # transforms.Normalize([meanR, meanG, meanB], [stdR, stdG, stdB]),
    # transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))

])

test_transformer_1 = transforms.Compose([
    transforms.Resize(512),
    transforms.RandomCrop(512),
    transforms.ToTensor(),
    # transforms.Normalize([meanR, meanG, meanB], [stdR, stdG, stdB]),

```

```
# transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))

])

train_ds.transform=train_transformer_1
test_ds.transform=test_transformer_1
```

▼ Data loader

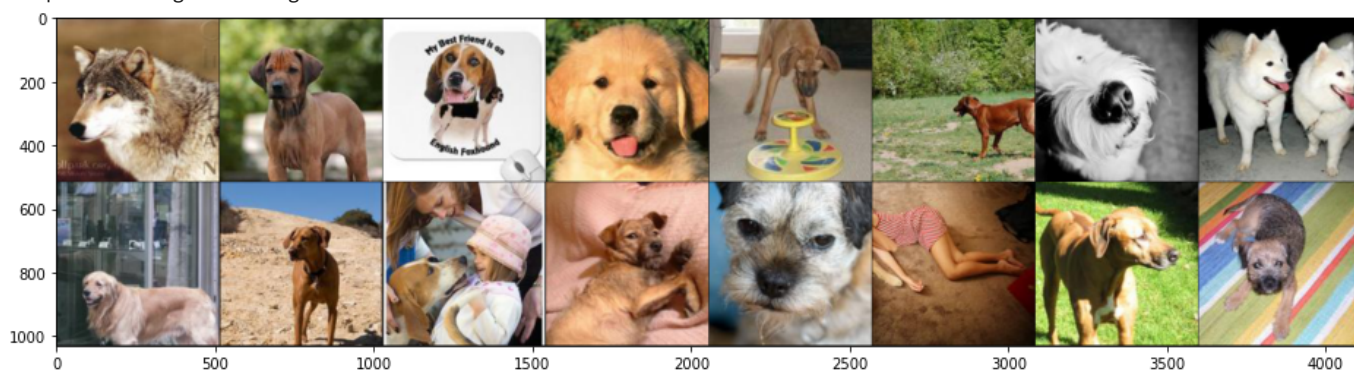
```
from torch.utils.data import DataLoader

train_dl = DataLoader(train_ds, batch_size = 32, shuffle=True, num_workers=4)
val_dl = DataLoader(val_ds, batch_size = 32, shuffle=True, num_workers=4)
test_dl = DataLoader(test_ds, batch_size = 32, shuffle=False, num_workers=4)
```

▼ Sample check

```
samples, labels = iter(train_dl).next()
plt.figure(figsize=(16,24))
grid_imgs = torchvision.utils.make_grid(samples[:24])
np_grid_imgs = grid_imgs.numpy()
# in tensor, image is (batch, width, height), so you have to transpose it to (width, height, batch) in numpy to show it.
plt.imshow(np.transpose(np_grid_imgs, (1,2,0)))
```

<matplotlib.image.AxesImage at 0x7fade135e978>



```
# extract a batch from training data
for x, y in train_dl:
    print(x.shape)
    print(y.shape)
    break

# extract a batch from val data
for x, y in val_dl:
    print(x.shape)
    print(y.shape)
    break
```

```
torch.Size([32, 3, 512, 512])
torch.Size([32])
torch.Size([32, 3, 512, 512])
torch.Size([32])
```

```
import collections

# get labels
y_train=[y for _,y in train_ds]

# count labels
```

```
counter_train=collections.Counter(y_train)
print(counter_train)
```

```
Counter({9: 664, 7: 664, 6: 660, 8: 659, 0: 659, 5: 658, 3: 652, 1: 650, 2: 645, 4: 406})
```

▼ Predefined functions

```
def train_val(model, params):
    # extract model parameters
    num_epochs=params["num_epochs"]
    loss_func=params["loss_func"]
    opt=params["optimizer"]
    train_dl=params["train_dl"]
    val_dl=params["val_dl"]
    sanity_check=params["sanity_check"]
    lr_scheduler=params["lr_scheduler"]
    path2weights=params["path2weights"]

    # history of loss values in each epoch
    loss_history={
        "train": [],
        "val": [],
    }

    # histroy of metric values in each epoch
    metric_history={
        "train": [],
        "val": [],
    }

    # a deep copy of weights for the best performing model
    best_model_wts = copy.deepcopy(model.state_dict())

    # initialize best loss to a large value
    best_loss=float('inf')

    # main loop
    for epoch in range(num_epochs):

        # get current learning rate
        current_lr=get_lr(opt)
        print('Epoch {}/{}. current lr={}'.format(epoch, num_epochs - 1, current_lr))

        # train model on training dataset
        model.train()
        train_loss, train_metric=loss_epoch(model,loss_func,train_dl,sanity_check,opt)

        # collect loss and metric for training dataset
        loss_history["train"].append(train_loss)
        metric_history["train"].append(train_metric)

        # evaluate model on validation dataset
        model.eval()
        with torch.no_grad():
            val_loss, val_metric=loss_epoch(model,loss_func,val_dl,sanity_check)

        # store best model
        if val_loss < best_loss:
            best_loss = val_loss
            best_model_wts = copy.deepcopy(model.state_dict())

        # store weights into a local file
        torch.save(model.state_dict(), path2weights)
        print("Copied best model weights!")
```

```

# collect loss and metric for validation dataset
loss_history["val"].append(val_loss)
metric_history["val"].append(val_metric)

# learning rate schedule
lr_scheduler.step()

print("train loss: %.6f, dev loss: %.6f, accuracy: %.2f" %(train_loss, val_loss, 100*val_metric))
print("-"*10)

# load best model weights
model.load_state_dict(best_model_wts)

return model, loss_history, metric_history

```

a helper function to compute the loss value and the performance metric for the entire dataset or an epoch.

```

# define device as a global variable
device = torch.device("cuda")

def loss_epoch(model, loss_func, dataset_dl, sanity_check=False, opt=None):
    running_loss=0.0
    running_metric=0.0
    len_data=len(dataset_dl.dataset)

    for xb, yb in dataset_dl:
        # move batch to device
        xb=xb.to(device)
        yb=yb.to(device)

        # get model output
        output=model(xb)

        # get loss per batch
        loss_b, metric_b=loss_batch(loss_func, output, yb, opt)

        # update running loss
        running_loss+=loss_b

        # update running metric
        if metric_b is not None:
            running_metric+=metric_b

        # break the loop in case of sanity check
        if sanity_check is True:
            break

    # average loss value
    loss=running_loss/float(len_data)

    # average metric value
    metric=running_metric/float(len_data)

    return loss, metric

```

A helper function to compute the loss value per batch of data:

```

def loss_batch(loss_func, output, target, opt=None):

```

```

    # get loss
    loss = loss_func(output, target)

    # get performance metric
    metric_b = metrics_batch(output, target)

    if opt is not None:
        opt.zero_grad()
        loss.backward()

```

```

    opt.step()

    return loss.item(), metric_b

```

A helper function to count the number of correct predictions per data batch:

```

def metrics_batch(output, target):
    # get output class
    pred = output.argmax(dim=1, keepdim=True)

    # compare output class with target class
    corrects=pred.eq(target.view_as(pred)).sum().item()
    return corrects

```

▼ Model_1

Resnet 18

```

from torchvision import models
import torch

# load model with random weights
model_1 = models.resnet18(pretrained=True)

```

```

print(model_1)

(layer2): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

        (relu): ReLU(inplace=True)
        (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )

```

```

# upload model to GPU
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

```

cuda:0

```

```

from torch import nn
# change the output layer
num_classes=10
num_fts = model_1.fc.in_features
model_1.fc = nn.Linear(num_fts, num_classes)

#device = torch.device("cuda:0")
model_1.to(device)

```

```

ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer3): Sequential(
    (0): BasicBlock(

```



```

(conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(downsample): Sequential(
  (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
  (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(1): BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)

```

```

# Even though the original image sizes are diff, we need to resize them to 224*224,
# the same size that the resnet18 model was trained at.
from torchsummary import summary
summary(model_1, input_size=(3, 224, 224))

```

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 64, 112, 112]	9,408
BatchNorm2d-2	[-1, 64, 112, 112]	128
ReLU-3	[-1, 64, 112, 112]	0
MaxPool2d-4	[-1, 64, 56, 56]	0
Conv2d-5	[-1, 64, 56, 56]	36,864
BatchNorm2d-6	[-1, 64, 56, 56]	128
ReLU-7	[-1, 64, 56, 56]	0
Conv2d-8	[-1, 64, 56, 56]	36,864
BatchNorm2d-9	[-1, 64, 56, 56]	128
ReLU-10	[-1, 64, 56, 56]	0
BasicBlock-11	[-1, 64, 56, 56]	0
Conv2d-12	[-1, 64, 56, 56]	36,864
BatchNorm2d-13	[-1, 64, 56, 56]	128
ReLU-14	[-1, 64, 56, 56]	0
Conv2d-15	[-1, 64, 56, 56]	36,864
BatchNorm2d-16	[-1, 64, 56, 56]	128
ReLU-17	[-1, 64, 56, 56]	0
BasicBlock-18	[-1, 64, 56, 56]	0
Conv2d-19	[-1, 128, 28, 28]	73,728
BatchNorm2d-20	[-1, 128, 28, 28]	256
ReLU-21	[-1, 128, 28, 28]	0
Conv2d-22	[-1, 128, 28, 28]	147,456
BatchNorm2d-23	[-1, 128, 28, 28]	256
Conv2d-24	[-1, 128, 28, 28]	8,192
BatchNorm2d-25	[-1, 128, 28, 28]	256
ReLU-26	[-1, 128, 28, 28]	0
BasicBlock-27	[-1, 128, 28, 28]	0
Conv2d-28	[-1, 128, 28, 28]	147,456
BatchNorm2d-29	[-1, 128, 28, 28]	256
ReLU-30	[-1, 128, 28, 28]	0
Conv2d-31	[-1, 128, 28, 28]	147,456
BatchNorm2d-32	[-1, 128, 28, 28]	256
ReLU-33	[-1, 128, 28, 28]	0
BasicBlock-34	[-1, 128, 28, 28]	0
Conv2d-35	[-1, 256, 14, 14]	294,912
BatchNorm2d-36	[-1, 256, 14, 14]	512
ReLU-37	[-1, 256, 14, 14]	0
Conv2d-38	[-1, 256, 14, 14]	589,824
BatchNorm2d-39	[-1, 256, 14, 14]	512
Conv2d-40	[-1, 256, 14, 14]	32,768
BatchNorm2d-41	[-1, 256, 14, 14]	512
ReLU-42	[-1, 256, 14, 14]	0
BasicBlock-43	[-1, 256, 14, 14]	0
Conv2d-44	[-1, 256, 14, 14]	589,824
BatchNorm2d-45	[-1, 256, 14, 14]	512
ReLU-46	[-1, 256, 14, 14]	0
Conv2d-47	[-1, 256, 14, 14]	589,824
BatchNorm2d-48	[-1, 256, 14, 14]	512
ReLU-49	[-1, 256, 14, 14]	0
BasicBlock-50	[-1, 256, 14, 14]	0
Conv2d-51	[-1, 512, 7, 7]	1,179,648

BatchNorm2d-52	[-1, 512, 7, 7]	1,024
ReLU-53	[-1, 512, 7, 7]	0
Conv2d-54	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-55	[-1, 512, 7, 7]	1,024
Conv2d-56	[-1, 512, 7, 7]	131,072
BatchNorm2d-57	[-1, 512, 7, 7]	1,024

▼ Training & Prediction

```
from torch import optim
opt = optim.Adam(model_1.parameters(), lr=1e-4)
```

```
# get learning rate
def get_lr(opt):
    for param_group in opt.param_groups:
        return param_group['lr']

current_lr=get_lr(opt)
print('current lr={}'.format(current_lr))
```

```
current lr=0.0001
```

```
from torch.optim.lr_scheduler import CosineAnnealingLR, StepLR

# define learning rate scheduler
# Several methods exist to adjust the learning rate. For a list of supported methods by PyTorch,
# please visit the following link: https://pytorch.org/docs/stable/optim.html.

#lr_scheduler = CosineAnnealingLR(opt,T_max=2,eta_min=1e-5)
lr_scheduler = StepLR(opt,step_size=1)
```

```
lrs=[]
for i in range(5):
    lr_scheduler.step()
    lr=get_lr(opt)
    print("epoch %s, lr: %.1e" %(i,lr))
    lrs.append(lr)
```

```
epoch 0, lr: 1.0e-05
epoch 1, lr: 1.0e-06
epoch 2, lr: 1.0e-07
epoch 3, lr: 1.0e-08
epoch 4, lr: 1.0e-09
/usr/local/lib/python3.6/dist-packages/torch/optim/lr_scheduler.py:136: UserWarning: Detected call of `lr_scheduler.step`
" https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
```

```
dict(counter_train)
```

```
{0: 753,
 1: 742,
 2: 737,
 3: 746,
 4: 464,
 5: 752,
 6: 754,
 7: 759,
 8: 754,
 9: 759}
```

```
from sortedcontainers import SortedDict

total_sum = sum(counter_train.values())

weights = SortedDict(dict(counter_train))
weights = nn.utils.weight_norm(weights, total_sum)
```

```
weights = np.array(list(weights.values()))
weights = weights/total_sum
weights = (torch.tensor(weights)).float().to(device)
weights

tensor([0.1043, 0.1028, 0.1021, 0.1033, 0.0643, 0.1042, 0.1044, 0.1051, 0.1044,
        0.1051], device='cuda:0')
```

```
import copy

loss_func = nn.CrossEntropyLoss(reduction="sum")

opt = optim.Adam(model_1.parameters(), lr=1e-4)
lr_scheduler = CosineAnnealingLR(opt,T_max=5,eta_min=1e-6)
#lr_scheduler = StepLR(opt,step_size=1)

params_train={
    "num_epochs": 5,
    "optimizer": opt,
    "loss_func": loss_func,
    "train_dl": train_dl,
    "val_dl": val_dl,
    "sanity_check": False,
    "lr_scheduler": lr_scheduler,
    "path2weights": folder / "resnet18_pretrained.pt",
}

# train and validate the model
model_1,loss_hist,metric_hist=train_val(model_1,params_train)

Epoch 0/4, current lr=0.0001
Copied best model weights!
train loss: 0.066807, dev loss: 0.343475, accuracy: 89.70
-----
Epoch 1/4, current lr=9.05463412215599e-05
train loss: 0.062225, dev loss: 0.387394, accuracy: 88.59
-----
Epoch 2/4, current lr=6.57963412215599e-05
Copied best model weights!
train loss: 0.030746, dev loss: 0.281684, accuracy: 91.41
-----
Epoch 3/4, current lr=3.52036587784401e-05
Copied best model weights!
train loss: 0.019883, dev loss: 0.267099, accuracy: 91.63
-----
Epoch 4/4, current lr=1.0453658778440105e-05
Copied best model weights!
train loss: 0.010560, dev loss: 0.255620, accuracy: 92.13
-----
```

```
from torch import nn
from torchvision import models
import torch

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

cuda:0
```

```
# load model
model_1 = models.resnet18(pretrained=True)
num_fts = model_1.fc.in_features
# change last layer
num_classes=10
model_1.fc = nn.Linear(num_fts, num_classes)
```

```
import torch

# load state_dict into model
path2weights=folder / "resnet18_pretrained.pt"
model_1.load_state_dict(torch.load(path2weights))
```

<All keys matched successfully>

```
test_dl.dataset.samples[0][0][78:]
```

'0.JPEG'

```
model_1.eval()
model_1.to(device)
fn_list = []
pred_list = []
for i, (x,fn) in enumerate(test_dl, 0):
    with torch.no_grad():
        x = x.to(device)
        output = model_1(x)
        pred = torch.argmax(output, dim=1)
        pred_list += [p.item() for p in pred]

for m,n in enumerate(test_dl.dataset.samples,0):
    fn_list += [n[0][78:]]

submission = pd.DataFrame({"file_names":fn_list, "target":pred_list})
submission.to_csv('preds_1.7.csv', index=False)
```

▼ Model_2

AlexNet

```
from torchvision import models
import torch
```

```
# load model with random weights
model_2 = models.alexnet(pretrained=True)
```

```
print(model_2)
```

```
AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
```

```

    )
)

# upload model to GPU
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

```
cuda:0
```

```
model_2.classifier[-1].out_features
```

```
1000
```

```

from torch import nn
# change the output layer
num_classes=10
num_fts = model_2.classifier[-1].out_features
model_2.fc = nn.Linear(num_fts, num_classes)

#device = torch.device("cuda:0")
model_2.to(device)

AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
  )
  (fc): Linear(in_features=1000, out_features=10, bias=True)
)

```

```

# Even though the original image sizes are 96*96, we need to resize them to 224*224,
# the same size that the resnet18 model was trained at.
from torchsummary import summary
summary(model_2, input_size=(3, 224, 224))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 55, 55]	23,296
ReLU-2	[-1, 64, 55, 55]	0
MaxPool2d-3	[-1, 64, 27, 27]	0
Conv2d-4	[-1, 192, 27, 27]	307,392
ReLU-5	[-1, 192, 27, 27]	0
MaxPool2d-6	[-1, 192, 13, 13]	0
Conv2d-7	[-1, 384, 13, 13]	663,936
ReLU-8	[-1, 384, 13, 13]	0
Conv2d-9	[-1, 256, 13, 13]	884,992
ReLU-10	[-1, 256, 13, 13]	0
Conv2d-11	[-1, 256, 13, 13]	590,080
ReLU-12	[-1, 256, 13, 13]	0
MaxPool2d-13	[-1, 256, 6, 6]	0

AdaptiveAvgPool2d-14	[-1, 256, 6, 6]	0
Dropout-15	[-1, 9216]	0
Linear-16	[-1, 4096]	37,752,832
ReLU-17	[-1, 4096]	0
Dropout-18	[-1, 4096]	0
Linear-19	[-1, 4096]	16,781,312
ReLU-20	[-1, 4096]	0
Linear-21	[-1, 1000]	4,097,000

```

=====
Total params: 61,100,840
Trainable params: 61,100,840
Non-trainable params: 0
-----

```

```

Input size (MB): 0.57
Forward/backward pass size (MB): 8.38
Params size (MB): 233.08
Estimated Total Size (MB): 242.03
-----

```

```

from torch import optim
opt = optim.Adam(model_2.parameters(), lr=1e-4)

```

```

# get learning rate
def get_lr(opt):
    for param_group in opt.param_groups:
        return param_group['lr']

current_lr=get_lr(opt)
print('current lr={}'.format(current_lr))

current_lr=0.0001

```

```

from torch.optim.lr_scheduler import CosineAnnealingLR

# define learning rate scheduler
# Several methods exist to adjust the learning rate. For a list of supported methods by PyTorch,
# please visit the following link: https://pytorch.org/docs/stable/optim.html.
lr_scheduler = CosineAnnealingLR(opt, T_max=2, eta_min=1e-5)

```

```

lrs=[]
for i in range(10):
    lr_scheduler.step()
    lr=get_lr(opt)
    print("epoch %s, lr: %.1e" %(i,lr))
    lrs.append(lr)

epoch 0, lr: 5.5e-05
epoch 1, lr: 1.0e-05
epoch 2, lr: 5.5e-05
epoch 3, lr: 1.0e-04
epoch 4, lr: 5.5e-05
epoch 5, lr: 1.0e-05
epoch 6, lr: 5.5e-05
epoch 7, lr: 1.0e-04
epoch 8, lr: 5.5e-05
epoch 9, lr: 1.0e-05
/usr/local/lib/python3.6/dist-packages/torch/optim/lr_scheduler.py:136: UserWarning: Detected call of `lr_scheduler.ste
https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)

```

```

import copy

loss_func = nn.CrossEntropyLoss(reduction="sum")
opt = optim.Adam(model_2.parameters(), lr=1e-4)
lr_scheduler = CosineAnnealingLR(opt, T_max=5, eta_min=1e-6)

params_train={
    "num_epochs": 4,
    "optimizer": opt,

```

```

"loss_func": loss_func,
"train_dl": train_dl,
"val_dl": val_dl,
"sanity_check": False,
"lr_scheduler": lr_scheduler,
"path2weights": folder / "alexnet_pretrained.pt",
}

# train and validate the model
model_2,loss_hist,metric_hist=train_val(model_2,params_train)

```

```

Epoch 0/3, current lr=0.0001
Copied best model weights!
train loss: 1.420489, dev loss: 0.610612, accuracy: 79.11
-----
Epoch 1/3, current lr=9.05463412215599e-05
Copied best model weights!
train loss: 0.571205, dev loss: 0.571547, accuracy: 81.61
-----
Epoch 2/3, current lr=6.57963412215599e-05
Copied best model weights!
train loss: 0.372884, dev loss: 0.515431, accuracy: 83.49
-----
Epoch 3/3, current lr=3.52036587784401e-05
Copied best model weights!
train loss: 0.238723, dev loss: 0.488585, accuracy: 85.15
-----

```

▼ Prediction

```

from torch import nn
from torchvision import models
import torch

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

```

cuda:0

```

```

# load model
model_2 = models.alexnet(pretrained=True)
num_fts = model_2.fc.in_features
# change last layer
num_classes=10
model_2.fc = nn.Linear(num_fts, num_classes)

```

```

import torch

# load state_dict into model
path2weights=folder / "alexnet_pretrained.pt"
model_2.load_state_dict(torch.load(path2weights))
    num_classes=10

test_dl.dataset.samples[0][0][78:]

    return modules[name]

model_2.eval()
model_2.to(device)
fn_list = []
pred_list = []
for i, (x,fn) in enumerate(test_dl, 0):
    with torch.no_grad():
        x = x.to(device)
        output = model_2(x)
        pred = torch.argmax(output, dim=1)
        pred_list += [p.item() for p in pred]

for m,n in enumerate(test_dl.dataset.samples,0):
    fn_list += [n[0][78:]]

submission = pd.DataFrame({"file_names":fn_list, "target":pred_list})
submission.to_csv('preds_2.1.csv', index=False)

```

▼ Model_3

Resnet 34

```

from torchvision import models
import torch

# load model with random weights
model_3= models.resnet34(pretrained=True)

```

Downloading: "<https://download.pytorch.org/models/resnet34-333f7ec4.pth>" to /root/.cache/torch/hub/checkpoints/resnet34
100% 83.3M/83.3M [00:07<00:00, 11.9MB/s]

```

print(model_3)
  (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
  (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(1): BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(2): BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(3): BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)

```



```

        (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
(4): BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(5): BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)

```

```

# upload model to GPU
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

```

cuda:0

```

```

from torch import nn
# change the output layer
num_classes=10
num_fts = model_3.fc.in_features
model_3.fc = nn.Linear(num_fts, num_classes)

#device = torch.device("cuda:0")
model_3.to(device)

    (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (2): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
  (3): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

```

```

(conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(reu): ReLU(inplace=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(4): BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (reu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
(5): BasicBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (reu): ReLU(inplace=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (reu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (reu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)

```

Even though the original image sizes are diff, we need to resize them to 224*224,
the same size that the resnet18 model was trained at.

```

from torchsummary import summary
summary(model_3, input_size=(3, 224, 224))

```

```

-----
  Conv2d-65      [-1, 256, 14, 14]      589,824
BatchNorm2d-66  [-1, 256, 14, 14]        512
  ReLU-67        [-1, 256, 14, 14]         0
  Conv2d-68      [-1, 256, 14, 14]      589,824
BatchNorm2d-69  [-1, 256, 14, 14]        512
  ReLU-70        [-1, 256, 14, 14]         0
  BasicBlock-71  [-1, 256, 14, 14]         0
  Conv2d-72      [-1, 256, 14, 14]      589,824
BatchNorm2d-73  [-1, 256, 14, 14]        512
  ReLU-74        [-1, 256, 14, 14]         0
  Conv2d-75      [-1, 256, 14, 14]      589,824
BatchNorm2d-76  [-1, 256, 14, 14]        512
  ReLU-77        [-1, 256, 14, 14]         0
  BasicBlock-78  [-1, 256, 14, 14]         0
  Conv2d-79      [-1, 256, 14, 14]      589,824
BatchNorm2d-80  [-1, 256, 14, 14]        512
  ReLU-81        [-1, 256, 14, 14]         0
  Conv2d-82      [-1, 256, 14, 14]      589,824
BatchNorm2d-83  [-1, 256, 14, 14]        512
  ReLU-84        [-1, 256, 14, 14]         0
  BasicBlock-85  [-1, 256, 14, 14]         0
  Conv2d-86      [-1, 256, 14, 14]      589,824
BatchNorm2d-87  [-1, 256, 14, 14]        512
  ReLU-88        [-1, 256, 14, 14]         0
  Conv2d-89      [-1, 256, 14, 14]      589,824
BatchNorm2d-90  [-1, 256, 14, 14]        512
  ReLU-91        [-1, 256, 14, 14]         0
  BasicBlock-92  [-1, 256, 14, 14]         0
  Conv2d-93      [-1, 256, 14, 14]      589,824
BatchNorm2d-94  [-1, 256, 14, 14]        512

```

ReLU-95	[-1, 256, 14, 14]	0
Conv2d-96	[-1, 256, 14, 14]	589,824
BatchNorm2d-97	[-1, 256, 14, 14]	512
ReLU-98	[-1, 256, 14, 14]	0
BasicBlock-99	[-1, 256, 14, 14]	0
Conv2d-100	[-1, 512, 7, 7]	1,179,648
BatchNorm2d-101	[-1, 512, 7, 7]	1,024
ReLU-102	[-1, 512, 7, 7]	0
Conv2d-103	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-104	[-1, 512, 7, 7]	1,024
Conv2d-105	[-1, 512, 7, 7]	131,072
BatchNorm2d-106	[-1, 512, 7, 7]	1,024
ReLU-107	[-1, 512, 7, 7]	0
BasicBlock-108	[-1, 512, 7, 7]	0
Conv2d-109	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-110	[-1, 512, 7, 7]	1,024
ReLU-111	[-1, 512, 7, 7]	0
Conv2d-112	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-113	[-1, 512, 7, 7]	1,024
ReLU-114	[-1, 512, 7, 7]	0
BasicBlock-115	[-1, 512, 7, 7]	0
Conv2d-116	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-117	[-1, 512, 7, 7]	1,024
ReLU-118	[-1, 512, 7, 7]	0
Conv2d-119	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-120	[-1, 512, 7, 7]	1,024
ReLU-121	[-1, 512, 7, 7]	0
BasicBlock-122	[-1, 512, 7, 7]	0
AdaptiveAvgPool2d-123	[-1, 512, 1, 1]	0

▼ Training & Prediction

```
from torch import optim
opt = optim.Adam(model_3.parameters(), lr=1e-4)
```

```
# get learning rate
def get_lr(opt):
    for param_group in opt.param_groups:
        return param_group['lr']

current_lr=get_lr(opt)
print('current lr={}'.format(current_lr))

current_lr=0.0001
```

```
from torch.optim.lr_scheduler import CosineAnnealingLR, StepLR

# define learning rate scheduler
# Several methods exist to adjust the learning rate. For a list of supported methods by PyTorch,
# please visit the following link: https://pytorch.org/docs/stable/optim.html.

lr_scheduler = CosineAnnealingLR(opt, T_max=2, eta_min=1e-5)
lr_scheduler = StepLR(opt, step_size=1)
```

```
lrs=[]
for i in range(5):
    lr_scheduler.step()
    lr=get_lr(opt)
    print("epoch %s, lr: %.1e" %(i,lr))
    lrs.append(lr)

epoch 0, lr: 1.0e-05
epoch 1, lr: 1.0e-06
epoch 2, lr: 1.0e-07
epoch 3, lr: 1.0e-08
epoch 4, lr: 1.0e-09
```

/usr/local/lib/python3.6/dist-packages/torch/optim/lr_scheduler.py:136: UserWarning: Detected call of `lr_scheduler.step` with parameters <https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate>, UserWarning)

◀ ▶

(counter_train)

```
Counter({0: 753,
         1: 742,
         2: 737,
         3: 746,
         4: 464,
         5: 752,
         6: 754,
         7: 759,
         8: 754,
         9: 759})
```

```
from sortedcontainers import SortedDict

total_sum = sum(counter_train.values())

weights = SortedDict(dict(counter_train))
weights = np.array(list(weights.values()))
weights = weights/total_sum
weights = (torch.tensor(weights)).float().to(device)
weights

tensor([0.1043, 0.1028, 0.1021, 0.1033, 0.0643, 0.1042, 0.1044, 0.1051, 0.1044,
        0.1051], device='cuda:0')
```

```
import copy

loss_func = nn.CrossEntropyLoss(reduction="sum")

opt = optim.Adam(model_3.parameters(), lr=1e-4)
lr_scheduler = CosineAnnealingLR(opt,T_max=5,eta_min=1e-6)
#lr_scheduler = StepLR(opt,step_size=1)

params_train={
    "num_epochs": 5,
    "optimizer": opt,
    "loss_func": loss_func,
    "train_dl": train_dl,
    "val_dl": val_dl,
    "sanity_check": False,
    "lr_scheduler": lr_scheduler,
    "path2weights": folder / "resnet34_pretrained.pt",
}

# train and validate the model
model_3,loss_hist,metric_hist=train_val(model_3,params_train)

Epoch 0/4, current lr=0.0001
Copied best model weights!
train loss: 0.462817, dev loss: 0.318257, accuracy: 89.58
-----
Epoch 1/4, current lr=9.05463412215599e-05
Copied best model weights!
train loss: 0.179780, dev loss: 0.266865, accuracy: 91.19
-----
Epoch 2/4, current lr=6.57963412215599e-05
Copied best model weights!
train loss: 0.069956, dev loss: 0.242825, accuracy: 92.47
-----
Epoch 3/4, current lr=3.52036587784401e-05
Copied best model weights!
train loss: 0.038639, dev loss: 0.223531, accuracy: 93.35
-----
Epoch 4/4, current lr=1.0453658778440105e-05
```

```
Copied best model weights!
train loss: 0.018329, dev loss: 0.214176, accuracy: 93.57
-----
```

```
from torch import nn
from torchvision import models
import torch

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)
```

```
cuda:0
```

```
# load model
model_3 = models.resnet34(pretrained=True)
num_fts = model_3.fc.in_features
# change last layer
num_classes=10
model_3.fc = nn.Linear(num_fts, num_classes)
```

```
import torch

# load state_dict into model
path2weights=folder / "resnet34_pretrained.pt"
model_3.load_state_dict(torch.load(path2weights))
```

```
<All keys matched successfully>
```

```
test_dl.dataset.samples[0][0][78:]
```

```
'0.JPEG'
```

```
model_3.eval()
model_3.to(device)
fn_list = []
pred_list = []
for i, (x,fn) in enumerate(test_dl, 0):
    with torch.no_grad():
        x = x.to(device)
        output = model_3(x)
        pred = torch.argmax(output, dim=1)
        pred_list += [p.item() for p in pred]

for m,n in enumerate(test_dl.dataset.samples,0):
    fn_list += [n[0][78:]]

submission = pd.DataFrame({"file_names":fn_list, "target":pred_list})
submission.to_csv('preds_3.1.csv', index=False)
```

▼ Model_4

Resnet 50

```
from torchvision import models
import torch

# load model with random weights
model_4= models.resnet50(pretrained=True)
```

```
print(model_4)
```

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): Bottleneck(
      (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (2): Bottleneck(
      (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
  (layer2): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): Bottleneck(
      (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
)
```

```
# upload model to GPU
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)
```

```
cuda:0
```

```
model_4
```

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
```

```

(bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace=True)
(maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
(layer1): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
  (2): Bottleneck(
    (conv1): Conv2d(256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
  )
)
(layer2): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)

```

```

from torch import nn
# change the output layer
num_classes=10
num_fts = model_4.fc.in_features
model_4.fc = nn.Linear(num_fts, num_classes)

```

```

#device = torch.device("cuda:0")
model_4.to(device)

```

```

(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace=True)
)
(2): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)

```

```

        (bn1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
    (3): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (4): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
    (5): Bottleneck(
      (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
    )
  )
  (layer4): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (downsample): Sequential(
        (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): Bottleneck(
      (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    )
  )

```

```

# Even though the original image sizes are diff, we need to resize them to 224*224,
# the same size that the resnet18 model was trained at.
from torchsummary import summary
summary(model_4, input_size=(3, 224, 224))

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 112, 112]	9,408
BatchNorm2d-2	[-1, 64, 112, 112]	128
ReLU-3	[-1, 64, 112, 112]	0
MaxPool2d-4	[-1, 64, 56, 56]	0
Conv2d-5	[-1, 64, 56, 56]	4,096
BatchNorm2d-6	[-1, 64, 56, 56]	128
ReLU-7	[-1, 64, 56, 56]	0
Conv2d-8	[-1, 64, 56, 56]	36,864
BatchNorm2d-9	[-1, 64, 56, 56]	128
ReLU-10	[-1, 64, 56, 56]	0
Conv2d-11	[-1, 256, 56, 56]	16,384
BatchNorm2d-12	[-1, 256, 56, 56]	512
Conv2d-13	[-1, 256, 56, 56]	16,384
BatchNorm2d-14	[-1, 256, 56, 56]	512
ReLU-15	[-1, 256, 56, 56]	0
Bottleneck-16	[-1, 256, 56, 56]	0

Conv2d-17	[-1, 64, 56, 56]	16,384
BatchNorm2d-18	[-1, 64, 56, 56]	128
ReLU-19	[-1, 64, 56, 56]	0
Conv2d-20	[-1, 64, 56, 56]	36,864
BatchNorm2d-21	[-1, 64, 56, 56]	128
ReLU-22	[-1, 64, 56, 56]	0
Conv2d-23	[-1, 256, 56, 56]	16,384
BatchNorm2d-24	[-1, 256, 56, 56]	512
ReLU-25	[-1, 256, 56, 56]	0
Bottleneck-26	[-1, 256, 56, 56]	0
Conv2d-27	[-1, 64, 56, 56]	16,384
BatchNorm2d-28	[-1, 64, 56, 56]	128
ReLU-29	[-1, 64, 56, 56]	0
Conv2d-30	[-1, 64, 56, 56]	36,864
BatchNorm2d-31	[-1, 64, 56, 56]	128
ReLU-32	[-1, 64, 56, 56]	0
Conv2d-33	[-1, 256, 56, 56]	16,384
BatchNorm2d-34	[-1, 256, 56, 56]	512
ReLU-35	[-1, 256, 56, 56]	0
Bottleneck-36	[-1, 256, 56, 56]	0
Conv2d-37	[-1, 128, 56, 56]	32,768
BatchNorm2d-38	[-1, 128, 56, 56]	256
ReLU-39	[-1, 128, 56, 56]	0
Conv2d-40	[-1, 128, 28, 28]	147,456
BatchNorm2d-41	[-1, 128, 28, 28]	256
ReLU-42	[-1, 128, 28, 28]	0
Conv2d-43	[-1, 512, 28, 28]	65,536
BatchNorm2d-44	[-1, 512, 28, 28]	1,024
Conv2d-45	[-1, 512, 28, 28]	131,072
BatchNorm2d-46	[-1, 512, 28, 28]	1,024
ReLU-47	[-1, 512, 28, 28]	0
Bottleneck-48	[-1, 512, 28, 28]	0
Conv2d-49	[-1, 128, 28, 28]	65,536
BatchNorm2d-50	[-1, 128, 28, 28]	256
ReLU-51	[-1, 128, 28, 28]	0
Conv2d-52	[-1, 128, 28, 28]	147,456
BatchNorm2d-53	[-1, 128, 28, 28]	256
ReLU-54	[-1, 128, 28, 28]	0
Conv2d-55	[-1, 512, 28, 28]	65,536
BatchNorm2d-56	[-1, 512, 28, 28]	1,024

▼ Training & Prediction

```
from torch import optim
opt = optim.Adam(model_4.parameters(), lr=1e-4)
```

```
# get learning rate
def get_lr(opt):
    for param_group in opt.param_groups:
        return param_group['lr']

current_lr=get_lr(opt)
print('current lr={}'.format(current_lr))

current_lr=0.0001
```

```
from torch.optim.lr_scheduler import CosineAnnealingLR, StepLR

# define learning rate scheduler
# Several methods exist to adjust the learning rate. For a list of supported methods by PyTorch,
# please visit the following link: https://pytorch.org/docs/stable/optim.html.


lr_scheduler = CosineAnnealingLR(opt,T_max=2,eta_min=1e-5)
lr_scheduler = StepLR(opt,step_size=1)
```

```
lrs=[]
for i in range(5):
    lr_scheduler.step()
    lr=get_lr(opt)
```

```
print("epoch %s, lr: %.1e" %(i,lr))
lrs.append(lr)
```

```
epoch 0, lr: 1.0e-05
epoch 1, lr: 1.0e-06
epoch 2, lr: 1.0e-07
epoch 3, lr: 1.0e-08
epoch 4, lr: 1.0e-09
```

```
/usr/local/lib/python3.6/dist-packages/torch/optim/lr_scheduler.py:136: UserWarning: Detected call of `lr_scheduler.step`
" https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
```



```
(counter_train)
```

```
Counter({0: 753,
         1: 742,
         2: 737,
         3: 746,
         4: 464,
         5: 752,
         6: 754,
         7: 759,
         8: 754,
         9: 759}))
```

```
from sortedcontainers import SortedDict
```

```
total_sum = sum(counter_train.values())
```

```
weights = SortedDict(dict(counter_train))
```

```
weights = np.array(list(weights.values()))
```

```
weights = weights/total_sum
```

```
weights = (torch.tensor(weights)).float().to(device)
```

```
weights
```

```
tensor([0.1043, 0.1028, 0.1021, 0.1033, 0.0643, 0.1042, 0.1044, 0.1051, 0.1044,
        0.1051], device='cuda:0')
```

```
import copy
```

```
loss_func = nn.CrossEntropyLoss(reduction="sum")
```

```
opt = optim.Adam(model_4.parameters(), lr=1e-4)
```

```
lr_scheduler = CosineAnnealingLR(opt,T_max=5,eta_min=1e-6)
```

```
#lr_scheduler = StepLR(opt,step_size=1)
```

```
params_train={
    "num_epochs": 5,
    "optimizer": opt,
    "loss_func": loss_func,
    "train_dl": train_dl,
    "val_dl": val_dl,
    "sanity_check": False,
    "lr_scheduler": lr_scheduler,
    "path2weights": folder / "resnet50_pretrained.pt",
}
```

```
# train and validate the model
```

```
model_4,loss_hist,metric_hist=train_val(model_4,params_train)
```

```
Epoch 0/4, current lr=0.0001
Copied best model weights!
train loss: 0.596847, dev loss: 0.329324, accuracy: 89.36
-----
Epoch 1/4, current lr=9.05463412215599e-05
Copied best model weights!
train loss: 0.293669, dev loss: 0.267223, accuracy: 91.08
-----
```

```

Epoch 2/4, current lr=6.57963412215599e-05
Copied best model weights!
train loss: 0.149720, dev loss: 0.266614, accuracy: 91.02
-----
Epoch 3/4, current lr=3.52036587784401e-05
Copied best model weights!
train loss: 0.067600, dev loss: 0.228830, accuracy: 93.85
-----
Epoch 4/4, current lr=1.0453658778440105e-05
Copied best model weights!
train loss: 0.038711, dev loss: 0.228743, accuracy: 92.74
-----

```

```

from torch import nn
from torchvision import models
import torch

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

```

cuda:0

```

```

# load model
model_4 = models.resnet50(pretrained=True)
num_fts = model_4.fc.in_features
# change last layer
num_classes=10
model_4.fc = nn.Linear(num_fts, num_classes)

```

```

import torch

# load state_dict into model
path2weights=folder / "resnet50_pretrained.pt"
model_4.load_state_dict(torch.load(path2weights))

```

```

<All keys matched successfully>

```

```

test_dl.dataset.samples[0][0][78:]

```

```

'0.JPEG'

```

```

model_4.eval()
model_4.to(device)
fn_list = []
pred_list = []
for i, (x,fn) in enumerate(test_dl, 0):
    with torch.no_grad():
        x = x.to(device)
        output = model_4(x)
        pred = torch.argmax(output, dim=1)
        pred_list += [p.item() for p in pred]

for m,n in enumerate(test_dl.dataset.samples,0):
    fn_list += [n[0][78:]]

submission = pd.DataFrame({"file_names":fn_list, "target":pred_list})
submission.to_csv('preds_4.3.csv', index=False)

```

▼ Model_5

Resnet 50 - fine tuning

```

from torchvision import models
import torch

# load model with random weights
model_5= models.resnet50(pretrained=True)

```

Downloading: "<https://download.pytorch.org/models/resnet50-19c8e357.pth>" to /root/.cache/torch/hub/checkpoints/resnet50
 100% 97.8M/97.8M [00:00<00:00, 285MB/s]

```
print(model_5)
```

```

(0): Bottleneck(
  (conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (downsample): Sequential(
    (0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
    (1): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  )
)
(1): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(2): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(3): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(4): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(5): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
)
(layer4): Sequential(

```

```
# upload model to GPU
```

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)
```

```
cuda:0
```

```
for param in model_5.layer4.parameters():
    param.requires_grad = True
```

```
from torch import nn
# change the output layer
num_classes=10
num_fts = model_5.fc.in_features
model_5.fc = nn.Linear(num_fts, num_classes)
```

```
#device = torch.device("cuda:0")
model_5.to(device)
```

```

    (relu): ReLU(inplace=True)
)
(2): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(3): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(4): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(5): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)

  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
)
(layer4): Sequential(
  (0): Bottleneck(
    (conv1): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
      (0): Conv2d(1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): Bottleneck(
    (conv1): Conv2d(2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

```

```
(conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv3): Conv2d(512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(2048, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
```

▼ Training & Prediction

```
from torch import optim
opt = optim.Adam(model_5.parameters(), lr=1e-4)
```

```
# get learning rate
def get_lr(opt):
    for param_group in opt.param_groups:
        return param_group['lr']

current_lr=get_lr(opt)
print('current lr={}'.format(current_lr))
```

```
current lr=0.0001
```

```
from torch.optim.lr_scheduler import CosineAnnealingLR, StepLR

# define learning rate scheduler
# Several methods exist to adjust the learning rate. For a list of supported methods by PyTorch,
# please visit the following link: https://pytorch.org/docs/stable/optim.html.

#lr_scheduler = CosineAnnealingLR(opt,T_max=2,eta_min=1e-5)
lr_scheduler = StepLR(opt,step_size=1)
```

```
lrs=[]
for i in range(5):
    lr_scheduler.step()
    lr=get_lr(opt)
    print("epoch %s, lr: %.1e" %(i,lr))
    lrs.append(lr)

epoch 0, lr: 1.0e-05
epoch 1, lr: 1.0e-06
epoch 2, lr: 1.0e-07
epoch 3, lr: 1.0e-08
epoch 4, lr: 1.0e-09
/usr/local/lib/python3.6/dist-packages/torch/optim/lr_scheduler.py:136: UserWarning: Detected call of `lr_scheduler.step`
" https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
```

```
(counter_train)
```

```
Counter({0: 753,
1: 742,
2: 737,
3: 746,
4: 464,
5: 752,
6: 754,
7: 759,
8: 754,
9: 759})
```

```
from sortedcontainers import SortedDict

total_sum = sum(counter_train.values())

weights = SortedDict(dict(counter_train))
weights = np.array(list(weights.values()))
weights = weights/total_sum
weights = (torch.tensor(weights)).float().to(device)
```

```
weights = torch.nn.Parameter(weights, requires_grad=True, device=device)
weights
```

```
tensor([0.1043, 0.1028, 0.1021, 0.1033, 0.0643, 0.1042, 0.1044, 0.1051, 0.1044,
        0.1051], device='cuda:0')
```

```
import copy

loss_func = nn.CrossEntropyLoss(reduction="sum")

opt = optim.Adam(model_5.parameters(), lr=1e-4)
lr_scheduler = CosineAnnealingLR(opt, T_max=5, eta_min=1e-6)
#lr_scheduler = StepLR(opt, step_size=1)

params_train={
    "num_epochs": 5,
    "optimizer": opt,
    "loss_func": loss_func,
    "train_dl": train_dl,
    "val_dl": val_dl,
    "sanity_check": False,
    "lr_scheduler": lr_scheduler,
    "path2weights": folder / "model_5.pt",
}

# train and validate the model
model_5, loss_hist, metric_hist = train_val(model_5, params_train)
```

```
Epoch 0/4, current lr=0.0001
Copied best model weights!
train loss: 0.591281, dev loss: 0.314889, accuracy: 90.75
-----
Epoch 1/4, current lr=9.05463412215599e-05
Copied best model weights!
train loss: 0.306739, dev loss: 0.300829, accuracy: 89.97
-----
Epoch 2/4, current lr=6.57963412215599e-05
Copied best model weights!
train loss: 0.143134, dev loss: 0.262663, accuracy: 91.86
-----
Epoch 3/4, current lr=3.52036587784401e-05
Copied best model weights!
train loss: 0.070000, dev loss: 0.216210, accuracy: 93.07
-----
Epoch 4/4, current lr=1.0453658778440105e-05
Copied best model weights!
train loss: 0.039831, dev loss: 0.193232, accuracy: 94.13
-----
```

```
from torch import nn
from torchvision import models
import torch

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)
```

```
cuda:0
```

```
# load model
model_5 = models.resnet50(pretrained=True)
num_fts = model_5.fc.in_features
# change last layer
num_classes=10
model_5.fc = nn.Linear(num_fts, num_classes)
```

```
import torch

# load state_dict into model
path2weights=folder / "model_5.pt"
```

```
model_5.load_state_dict(torch.load(path2weights))
```

```
<All keys matched successfully>
```

```
test_dl.dataset.samples[0][0][78:]
```

```
'0.JPEG'
```

```
model_5.eval()
model_5.to(device)
fn_list = []
pred_list = []
for i, (x,fn) in enumerate(test_dl, 0):
    with torch.no_grad():
        x = x.to(device)
        output = model_5(x)
        pred = torch.argmax(output, dim=1)
        pred_list += [p.item() for p in pred]

for m,n in enumerate(test_dl.dataset.samples,0):
    fn_list += [n[0][78:]]

submission = pd.DataFrame({"file_names":fn_list, "target":pred_list})
#submission.to_csv('submission.csv', index=False)
```

```
submission.to_csv('pred_5.csv', index=False)
```

▼ Model_6

Resnet 152

```
from torchvision import models
import torch

# load model with random weights
model_6= models.resnet152(pretrained=True)
```

Downloading: "<https://download.pytorch.org/models/resnet152-b121ed2d.pth>" to /root/.cache/torch/hub/checkpoints/resnet152-b121ed2d.pth
100% 230M/230M [00:06<00:00, 38.6MB/s]

```
print(model_6)
...
    (relu): ReLU(inplace=True)
  )
(18): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(19): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
```



```

)
(20): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(21): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(22): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(23): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(24): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

# upload model to GPU
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

```

cuda:0

```

```

from torch import nn
# change the output layer
num_classes=10
num_fters = model_6.fc.in_features
model_6.fc = nn.Linear(num_fters, num_classes)

#device = torch.device("cuda:0")
model_6.to(device)

  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(28): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(29): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)

```

```

        (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu): ReLU(inplace=True)
    )
(30): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)

  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(31): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(32): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(33): Bottleneck(
  (conv1): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
)
(34): Bottleneck(

```

▼ Training & Prediction

```

from torch import optim
opt = optim.Adam(model_6.parameters(), lr=1e-4)

```

```

# get learning rate
def get_lr(opt):
    for param_group in opt.param_groups:
        return param_group['lr']

current_lr=get_lr(opt)
print('current lr={}'.format(current_lr))

```

```
current_lr=0.0001
```

```

from torch.optim.lr_scheduler import CosineAnnealingLR, StepLR

# define learning rate scheduler
# Several methods exist to adjust the learning rate. For a list of supported methods by PyTorch,
# please visit the following link: https://pytorch.org/docs/stable/optim.html.

lr_scheduler = CosineAnnealingLR(opt, T_max=2, eta_min=1e-5)
lr_scheduler = StepLR(opt, step_size=1)

```

```

lrs=[]
for i in range(5):

```

```

lr_scheduler.step()
lr=get_lr(opt)
print("epoch %s, lr: %.1e" %(i,lr))
lrs.append(lr)

```

```

epoch 0, lr: 1.0e-05
epoch 1, lr: 1.0e-06
epoch 2, lr: 1.0e-07
epoch 3, lr: 1.0e-08
epoch 4, lr: 1.0e-09

```

```

/usr/local/lib/python3.6/dist-packages/torch/optim/lr_scheduler.py:136: UserWarning: Detected call of `lr_scheduler.step`
with argument `step_size` which has no effect on the scheduler.
  """https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)

```

◀ ▶

```

(counter_train)

```

```

Counter({0: 753,
         1: 742,
         2: 737,
         3: 746,
         4: 464,
         5: 752,
         6: 754,
         7: 759,
         8: 754,
         9: 759})

```

```

from sortedcontainers import SortedDict

```

```

total_sum = sum(counter_train.values())

```

```

weights = SortedDict(dict(counter_train))

```

```

weights = np.array(list(weights.values()))

```

```

weights = weights/total_sum

```

```

weights = (torch.tensor(weights)).float().to(device)

```

```

weights

```

```

tensor([0.1043, 0.1028, 0.1021, 0.1033, 0.0643, 0.1042, 0.1044, 0.1051, 0.1044,
        0.1051], device='cuda:0')

```

```

import copy

```

```

loss_func = nn.CrossEntropyLoss(reduction="sum")

```

```

opt = optim.Adam(model_6.parameters(), lr=1e-4)

```

```

lr_scheduler = CosineAnnealingLR(opt,T_max=5,eta_min=1e-6)

```

```

#lr_scheduler = StepLR(opt,step_size=1)

```

```

params_train={

```

```

    "num_epochs": 5,

```

```

    "optimizer": opt,

```

```

    "loss_func": loss_func,

```

```

    "train_dl": train_dl,

```

```

    "val_dl": val_dl,

```

```

    "sanity_check": False,

```

```

    "lr_scheduler": lr_scheduler,

```

```

    "path2weights": folder / "model_6.pt",

```

```

}

```

```

# train and validate the model

```

```

model_6,loss_hist,metric_hist=train_val(model_6,params_train)

```

```

Epoch 0/4, current lr=0.0001

```

```

Copied best model weights!

```

```

train loss: 0.587427, dev loss: 0.422705, accuracy: 85.32

```

```

-----

```

```

Epoch 1/4, current lr=9.05463412215599e-05

```

```

Copied best model weights!

```

```

train loss: 0.303875, dev loss: 0.379381, accuracy: 87.87
-----
Epoch 2/4, current lr=6.57963412215599e-05
Copied best model weights!
train loss: 0.171443, dev loss: 0.330739, accuracy: 90.03
-----
Epoch 3/4, current lr=3.52036587784401e-05
Copied best model weights!
train loss: 0.063638, dev loss: 0.243718, accuracy: 92.41
-----
Epoch 4/4, current lr=1.0453658778440105e-05
Copied best model weights!
train loss: 0.032323, dev loss: 0.208538, accuracy: 93.68
-----

```

```

from torch import nn
from torchvision import models
import torch

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

cuda:0

```

```

# load model
model_6 = models.resnet152(pretrained=True)
num_fts = model_6.fc.in_features
# change last layer
num_classes=10
model_6.fc = nn.Linear(num_fts, num_classes)

```

```

import torch

# load state_dict into model
path2weights=folder / "model_6.pt"
model_6.load_state_dict(torch.load(path2weights))

```

```

<All keys matched successfully>

```

```

test_dl.dataset.samples[0][0][78:]

'0.JPEG'

```

```

model_6.eval()
model_6.to(device)
fn_list = []
pred_list = []
for i, (x,fn) in enumerate(test_dl, 0):
    with torch.no_grad():
        x = x.to(device)
        output = model_6(x)
        pred = torch.argmax(output, dim=1)
        pred_list += [p.item() for p in pred]

for m,n in enumerate(test_dl.dataset.samples,0):
    fn_list += [n[0][78:]]

submission = pd.DataFrame({"file_names":fn_list, "target":pred_list})
submission.to_csv('preds_6.csv', index=False)

```

▼ Model_7

Densenet 161

```
from torchvision import models
import torch

# load model with random weights
model_7= models.densenet161(pretrained=True)

print(model_7)

    (relu1): ReLU(inplace=True)
    (conv1): Conv2d(96, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu2): ReLU(inplace=True)
    (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  )
  (denselayer2): _DenseLayer(
    (norm1): BatchNorm2d(144, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu1): ReLU(inplace=True)
    (conv1): Conv2d(144, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu2): ReLU(inplace=True)
    (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  )
  (denselayer3): _DenseLayer(
    (norm1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu1): ReLU(inplace=True)
    (conv1): Conv2d(192, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu2): ReLU(inplace=True)
    (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  )
  (denselayer4): _DenseLayer(
    (norm1): BatchNorm2d(240, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu1): ReLU(inplace=True)
    (conv1): Conv2d(240, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu2): ReLU(inplace=True)
    (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  )
  (denselayer5): _DenseLayer(
    (norm1): BatchNorm2d(288, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu1): ReLU(inplace=True)
    (conv1): Conv2d(288, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu2): ReLU(inplace=True)
    (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  )
  (denselayer6): _DenseLayer(
    (norm1): BatchNorm2d(336, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu1): ReLU(inplace=True)
    (conv1): Conv2d(336, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu2): ReLU(inplace=True)
    (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  )
  )
  (transition1): _Transition(
    (norm): BatchNorm2d(384, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv): Conv2d(384, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (pool): AvgPool2d(kernel_size=2, stride=2, padding=0)
  )
  (denseblock2): _DenseBlock(
    (denselayer1): _DenseLayer(
      (norm1): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu1): ReLU(inplace=True)
      (conv1): Conv2d(192, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu2): ReLU(inplace=True)
      (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
  )

# upload model to GPU
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)
```

cuda:0

```
from torch import nn
# change the output layer
num_classes=10

model_7.classifier.out_features = num_classes

#device = torch.device("cuda:0")
model_7.to(device)

        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
    (denselayer9): _DenseLayer(
        (norm1): BatchNorm2d(1440, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu1): ReLU(inplace=True)
        (conv1): Conv2d(1440, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
    (denselayer10): _DenseLayer(
        (norm1): BatchNorm2d(1488, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu1): ReLU(inplace=True)
        (conv1): Conv2d(1488, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
    (denselayer11): _DenseLayer(
        (norm1): BatchNorm2d(1536, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu1): ReLU(inplace=True)
        (conv1): Conv2d(1536, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
    (denselayer12): _DenseLayer(
        (norm1): BatchNorm2d(1584, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu1): ReLU(inplace=True)
        (conv1): Conv2d(1584, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
    (denselayer13): _DenseLayer(
        (norm1): BatchNorm2d(1632, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu1): ReLU(inplace=True)
        (conv1): Conv2d(1632, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
    (denselayer14): _DenseLayer(
        (norm1): BatchNorm2d(1680, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu1): ReLU(inplace=True)
        (conv1): Conv2d(1680, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
    (denselayer15): _DenseLayer(
        (norm1): BatchNorm2d(1728, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu1): ReLU(inplace=True)
        (conv1): Conv2d(1728, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (relu2): ReLU(inplace=True)
        (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    )
    (denselayer16): DenseLayer(
```

▼ Training & Prediction

```
from torch import optim
opt = optim.Adam(model_7.parameters(), lr=1e-4)
```

```
# get learning rate
def get_lr(opt):
    for param_group in opt.param_groups:
        return param_group['lr']
```

```
current_lr=get_lr(opt)
print('current lr={}'.format(current_lr))
```

```
current_lr=0.0001
```

```
from torch.optim.lr_scheduler import CosineAnnealingLR, StepLR
```

```
# define learning rate scheduler
# Several methods exist to adjust the learning rate. For a list of supported methods by PyTorch,
# please visit the following link: https://pytorch.org/docs/stable/optim.html.
```

```
#lr_scheduler = CosineAnnealingLR(opt,T_max=2,eta_min=1e-5)
lr_scheduler = StepLR(opt,step_size=1)
```

```
lrs=[]
for i in range(5):
    lr_scheduler.step()
    lr=get_lr(opt)
    print("epoch %s, lr: %.1e" % (i,lr))
    lrs.append(lr)
```

```
epoch 0, lr: 1.0e-05
epoch 1, lr: 1.0e-06
epoch 2, lr: 1.0e-07
epoch 3, lr: 1.0e-08
epoch 4, lr: 1.0e-09
```

```
/usr/local/lib/python3.6/dist-packages/torch/optim/lr_scheduler.py:136: UserWarning: Detected call of `lr_scheduler.step`
" https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate", UserWarning)
```

```
(counter_train)
```

```
Counter({0: 753,
1: 742,
2: 737,
3: 746,
4: 464,
5: 752,
6: 754,
7: 759,
8: 754,
9: 759})
```

```
from sortedcontainers import SortedDict
```

```
total_sum = sum(counter_train.values())
```

```
weights = SortedDict(dict(counter_train))
weights = np.array(list(weights.values()))
weights = np.around(weights/total_sum,decimals=2)
weights = (torch.tensor(weights)).float().to(device)
weights
```

```
tensor([0.1000, 0.1000, 0.1000, 0.1000, 0.0600, 0.1000, 0.1000, 0.1100, 0.1000,
0.1100], device='cuda:0')
```

```
import copy
```

```

loss_func = nn.CrossEntropyLoss(reduction="sum")

opt = optim.Adam(model_7.parameters(), lr=1e-4)
lr_scheduler = CosineAnnealingLR(opt, T_max=5, eta_min=1e-6)
#lr_scheduler = StepLR(opt, step_size=1)

params_train={
    "num_epochs": 5,
    "optimizer": opt,
    "loss_func": loss_func,
    "train_dl": train_dl,
    "val_dl": val_dl,
    "sanity_check": False,
    "lr_scheduler": lr_scheduler,
    "path2weights": folder / "model_7.pt",
}

# train and validate the model
model_7, loss_hist, metric_hist = train_val(model_7, params_train)

```

```

Epoch 0/4, current lr=0.0001
Copied best model weights!
train loss: 0.757316, dev loss: 0.309049, accuracy: 89.92
-----
Epoch 1/4, current lr=9.05463412215599e-05
Copied best model weights!
train loss: 0.217857, dev loss: 0.261615, accuracy: 91.19
-----
Epoch 2/4, current lr=6.57963412215599e-05
Copied best model weights!
train loss: 0.092043, dev loss: 0.246794, accuracy: 93.91
-----
Epoch 3/4, current lr=3.52036587784401e-05
Copied best model weights!
train loss: 0.041339, dev loss: 0.224900, accuracy: 93.68
-----
Epoch 4/4, current lr=1.0453658778440105e-05
Copied best model weights!
train loss: 0.020577, dev loss: 0.196145, accuracy: 94.85
-----

```

```

from torch import nn
from torchvision import models
import torch

device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

```

cuda:0

```

```

# load model
model_7 = models.densenet161(pretrained=True)

from torch import nn
# change the output layer
num_classes=10

model_7.classifier.out_features = num_classes

#device = torch.device("cuda:0")
model_7.to(device)

(conv1): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
(norm2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu2): ReLU(inplace=True)
(conv2): Conv2d(128, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)
(denselayer8): _DenseLayer(
  (norm1): BatchNorm2d(1392, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(1392, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
)

```



```

(norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu2): ReLU(inplace=True)
(conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)
(denselayer9): _DenseLayer(
  (norm1): BatchNorm2d(1440, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(1440, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu2): ReLU(inplace=True)
  (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)
(denselayer10): _DenseLayer(
  (norm1): BatchNorm2d(1488, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(1488, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu2): ReLU(inplace=True)
  (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)
(denselayer11): _DenseLayer(
  (norm1): BatchNorm2d(1536, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(1536, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu2): ReLU(inplace=True)
  (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)
(denselayer12): _DenseLayer(
  (norm1): BatchNorm2d(1584, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(1584, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu2): ReLU(inplace=True)
  (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)
(denselayer13): _DenseLayer(
  (norm1): BatchNorm2d(1632, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(1632, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu2): ReLU(inplace=True)
  (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)
(denselayer14): _DenseLayer(
  (norm1): BatchNorm2d(1680, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu1): ReLU(inplace=True)
  (conv1): Conv2d(1680, 192, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (norm2): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu2): ReLU(inplace=True)
  (conv2): Conv2d(192, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
)

```

```
import torch
```

```

# load state_dict into model
path2weights=folder / "model_7.pt"
model_7.load_state_dict(torch.load(path2weights))

```

```
<All keys matched successfully>
```

```
test_dl.dataset.samples[0][0][78:]
```

```
'0.JPEG'
```

```

model_7.eval()
model_7.to(device)
fn_list = []
pred_list = []
for i, (x,fn) in enumerate(test_dl, 0):
    with torch.no_grad():
        x = x.to(device)
        output = model_7(x)
        pred = torch.argmax(output, dim=1)

```

```

pred = torch.argmax(output, dim=1)
pred_list += [p.item() for p in pred]

for m,n in enumerate(test_dl.dataset.samples,0):
    fn_list += [n[0][78:]]

submission = pd.DataFrame({"file_names":fn_list, "target":pred_list})
submission.to_csv('preds_7.csv', index=False)

```

▼ Model_8

```

from torchvision import models
import torch

# load model with random weights
model_8= models.vgg19_bn(pretrained=True)

```

Downloading: "https://download.pytorch.org/models/vgg19_bn-c79401a0.pth" to /root/.cache/torch/hub/checkpoints/vgg19_bn-100%
548M/548M [00:12<00:00, 45.7MB/s]

```

print(model_8)
(4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(5): ReLU(inplace=True)
(6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(7): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(9): ReLU(inplace=True)
(10): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(12): ReLU(inplace=True)
(13): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(14): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(16): ReLU(inplace=True)
(17): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(19): ReLU(inplace=True)
(20): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(21): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(22): ReLU(inplace=True)
(23): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(25): ReLU(inplace=True)
(26): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(27): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(28): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(32): ReLU(inplace=True)
(33): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(34): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(35): ReLU(inplace=True)
(36): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(37): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(38): ReLU(inplace=True)
(39): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(40): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(41): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(42): ReLU(inplace=True)
(43): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(44): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(45): ReLU(inplace=True)
(46): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(47): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(48): ReLU(inplace=True)
(49): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(50): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(51): ReLU(inplace=True)
(52): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace=True)
  (5): Dropout(p=0.5, inplace=False)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)

```

```

# upload model to GPU
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(device)

```

```

cuda:0

```

```

#for param in model_8.classifier.parameters():
# param.requires_grad_ = True

```

```

model_8.classifier[-1].out_features

```

```

1000

```

```

from torch import nn
# change the output layer
num_classes=10
model_8.classifier[-1].out_features = num_classes

```

```

#device = torch.device("cuda:0")
model_8.to(device)

```

```

(4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(5): ReLU(inplace=True)
(6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(7): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(9): ReLU(inplace=True)
(10): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(12): ReLU(inplace=True)
(13): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(14): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(16): ReLU(inplace=True)
(17): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(19): ReLU(inplace=True)
(20): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(21): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(22): ReLU(inplace=True)
(23): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(25): ReLU(inplace=True)
(26): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(27): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(28): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(32): ReLU(inplace=True)
(33): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(34): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(35): ReLU(inplace=True)
(36): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(37): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)

```

```

(38): ReLU(inplace=True)
(39): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(40): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(41): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(42): ReLU(inplace=True)
(43): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(44): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(45): ReLU(inplace=True)
(46): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(47): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(48): ReLU(inplace=True)
(49): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(50): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(51): ReLU(inplace=True)
(52): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace=True)
  (5): Dropout(p=0.5, inplace=False)
  (6): Linear(in_features=4096, out_features=10, bias=True)
)

```

▼ Training & Prediction

```

from torch import optim
opt = optim.Adam(model_8.parameters(), lr=1e-4)

```

```

# get learning rate
def get_lr(opt):
    for param_group in opt.param_groups:
        return param_group['lr']

current_lr=get_lr(opt)
print('current lr={}'.format(current_lr))

```

```
current_lr=0.0001
```

```

from torch.optim.lr_scheduler import CosineAnnealingLR, StepLR

# define learning rate scheduler
# Several methods exist to adjust the learning rate. For a list of supported methods by PyTorch,
# please visit the following link: https://pytorch.org/docs/stable/optim.html.

#lr_scheduler = CosineAnnealingLR(opt,T_max=2,eta_min=1e-5)
lr_scheduler = StepLR(opt,step_size=1)

```

```

lrs=[]
for i in range(5):
    lr_scheduler.step()
    lr=get_lr(opt)
    print("epoch %s, lr: %.1e" %(i,lr))
    lrs.append(lr)

```

```

epoch 0, lr: 1.0e-05
epoch 1, lr: 1.0e-06
epoch 2, lr: 1.0e-07
epoch 3, lr: 1.0e-08
epoch 4, lr: 1.0e-09

```

```

/usr/local/lib/python3.6/dist-packages/torch/optim/lr_scheduler.py:136: UserWarning: Detected call of `lr_scheduler.step`
with argument `step_size` that has been deprecated since version 3.0.0. Please see https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate for details.

```

```
(counter_train)
```

```

from sortedcontainers import SortedDict

total_sum = sum(counter_train.values())

weights = SortedDict(dict(counter_train))
weights = np.array(list(weights.values()))
weights = np.around(weights/total_sum,decimals=2)
weights = (torch.tensor(weights)).float().to(device)
weights

```

```

import copy

loss_func = nn.CrossEntropyLoss(reduction="sum")

opt = optim.Adam(model_8.parameters(), lr=1e-4)
lr_scheduler = CosineAnnealingLR(opt,T_max=5,eta_min=1e-6)
#lr_scheduler = StepLR(opt,step_size=1)

params_train={
    "num_epochs": 5,
    "optimizer": opt,
    "loss_func": loss_func,
    "train_dl": train_dl,
    "val_dl": val_dl,
    "sanity_check": False,
    "lr_scheduler": lr_scheduler,
    "path2weights": folder / "model_8.pt",
}

# train and validate the model
model_8,loss_hist,metric_hist=train_val(model_8,params_train)

```

```

test_dl.dataset.samples[0][0][78:]

```

```

model_8.eval()
model_8.to(device)
fn_list = []
pred_list = []
for i, (x,fn) in enumerate(test_dl, 0):
    with torch.no_grad():
        x = x.to(device)
        output = model_8(x)
        pred = torch.argmax(output, dim=1)
        pred_list += [p.item() for p in pred]

for m,n in enumerate(test_dl.dataset.samples,0):
    fn_list += [n[0][78:]]

submission = pd.DataFrame({"file_names":fn_list, "target":pred_list})
submission.to_csv('preds_8.csv', index=False)

```

